

● Hub71 Application · Abu Dhabi · 2026

EVE

*The ambient intelligence that thinks about your life — so
you don't have to.*

- Ambient AI Agent
- Voice-First
- Graph Memory
- Mobile-Native
- B2C SaaS
- Pre-Seed

The digital world demands more attention than any human has.

"I spend the first two hours of every morning just figuring out what I'm supposed to do — before I do any actual work."

— Every professional we interviewed

2.5h

Lost daily to inbox triage, calendar management & context-switching

147

Average emails received per day by a professional in 2026

73%

Of professionals feel overwhelmed by digital communication daily

OpenClaw/Moltbot proved the demand: **145,000 GitHub stars in three weeks**. But it was desktop-only, required a self-hosted server, stored credentials in plaintext, and cost users \$300+/month to run. The problem is real. Nobody has solved it correctly — especially not on mobile.

EVE doesn't wait to be asked. She already knows.

EVE is the world's first ambient AI agent designed for mobile — a persistent intelligence that lives on your phone, monitors your digital life server-side, and proactively reaches out when something needs your attention.

She speaks. She listens. She remembers — not as raw transcripts, but through a graph-structured knowledge model that gets sharper the longer you use her, without ever slowing down or costing more.

Proactive — she contacts you first

When an urgent email arrives or a meeting approaches, EVE taps you on the shoulder. Not a notification. A conversation.

Voice-native — natural conversation

Sub-400ms response latency. Barge-in support. Natural hesitations and rhythm. The uncanny valley of AI voice, solved.

Graph memory — gets smarter over time

A versioned knowledge graph, not a flat vector store. Precise, auditable, and bounded — so a power user after 8 months costs the same as a new user.

EVE IN ACTION — 8:03 PM

- 1

EVE detects a pattern
Interview tomorrow at 9am. Project Atlas — no activity in 7 days. Correlation flagged.
- 2

EVE initiates conversation
"Hey — big day tomorrow. How's Project Atlas going? I had a few feature ideas I wanted to run by you."
- 3

Live voice brainstorm
A full conversational session — EVE surfaces ideas, asks questions, builds on the user's answers in real time.
- 4

Memory updates silently
New decisions, new facts — compressed into the graph for next time. Zero extra cost, zero user friction.

Three layers. One seamless experience.

Context Compiler

Core IP

A deterministic Go service that assembles a fixed-size context package — Tier 1 (identity core, pre-cached), Tier 2 (graph-retrieved situational context), Tier 3 (live turn). A power user with 8 months of history gets the same prompt size and latency as a new user. More history → better curation, not bigger cost.

Graph Memory Engine

Moat

A versioned knowledge graph in PostgreSQL — entities, relationships, temporal edges. 80% of queries resolve via pure graph traversal (zero embedding cost). Vector search as fallback for fuzzy queries only. Memory is transparent — users see and edit what EVE knows.

Voice Pipeline

V2

LiveKit WebRTC transport + ElevenLabs Turbo TTS with streaming chunk synthesis. Sentence-boundary buffering starts audio within 400ms. Silero VAD on-device for instant barge-in. OpenAI Realtime API for voice-to-voice mode when latency is the priority.

MODEL TIERING STRATEGY

Triage polling (15min)

Haiku — \$0.001/call

Morning briefing

Sonnet — \$0.03/call

Voice conversation

Sonnet — streaming

Proactive brainstorm

Opus — rare, high-value

STACK

MOBILE	React Native + Expo (iOS + Android)
BACKEND	Go + Fiber — goroutines for 10K+ concurrent WS
JOBs	Asynq (Redis-backed) — scheduled agent tasks
DATA	PostgreSQL + pgvector + sqlc
INFRA	Fly.io edge — sub-100ms from MENA

The quadrant nobody is building.

Product	Mobile	Proactive	Voice	Graph memory	Managed	<\$20/mo
EVE	✓	✓	✓	✓	✓	✓
OpenClaw	✗	~	✗	✗	✗	✗
Hermes Agent	~	~	✗	✗	✗	✗
ChatGPT	~	✗	~	✗	✓	~
Superhuman	✓	✗	✗	✗	✓	✗
Gemini Live	✓	✗	✓	✗	✓	✓

The **proactive + mobile + voice + graph memory + managed** combination is completely unoccupied. Every competitor owns at most two of these attributes.

vs OpenClaw / Hermes

Both require self-hosting, BYOK, and technical setup. Their "mobile" story is a remote viewer for a server the user runs themselves. EVE requires zero setup and bundles API costs in the subscription.

vs ChatGPT / Gemini

Reactive by design — they wait to be asked. No proactive interruptions, no graph memory that persists across sessions in a structured form, no ambient phone presence. They're tools, not agents.

vs Big Tech's roadmap

Google and Apple will ship something eventually. EVE's moat is the mobile-first African and MENA market — 2 billion users Big Tech has structurally ignored — plus the Context Compiler and graph memory architecture that can't be replicated in a "feature drop."

A category being born. A region being ignored.



\$370M invested in ambient agent infrastructure in 2025–26

The deployment phase is 2026–2028. The window to establish category position is open — but it is closing. EVE is built to move now.

STRATEGIC BEACHHEAD — MENA + AFRICA

500M+ smartphone users across MENA and Sub-Saharan Africa. Mobile-first by necessity. WhatsApp as the communication OS. Massive underserved demand for AI productivity tools that actually work on Android.

500M+

MENA + Africa smartphone users

\$32B

MENA digital economy by 2030

Abu Dhabi as global AI hub

UAE National AI Strategy 2031, G42 ecosystem, Mubadala’s global tech portfolio — Abu Dhabi is the ideal launchpad for an AI company targeting emerging mobile markets. Hub71’s network provides what no other accelerator in the world can: direct access to the world’s fastest-growing tech capital.

OpenClaw proved the demand signal

145K GitHub stars. \$16M in fraudulent crypto tokens exploiting the hype. Three rebrands in one week. The demand is beyond dispute. No clean mobile product has captured it.

~80% gross margin.
From day one.

STARTER

\$0 / 7 days

→ Gmail + Calendar

→ 1 briefing/day

→ 5 reply approvals

→ No credit card

Most popular

PRO

\$12 / month

→ Unlimited approvals

→ Custom briefing time

→ Proactive EVE interrupts (V2)

→ Voice conversations (V2)

→ Transparent memory graph

POWER

\$19 / month

→ Everything in Pro

→ Outlook + Slack

→ Bring your own API key

→ Multiple briefing slots

→ Voice-to-voice mode

≤\$0.08

Max LLM cost per user per day — enforced at infra level

~80%

Gross margin on LLM costs at Pro tier. \$2.40/month in API costs.

\$36K

MRR target at month 12 — 3,000 users, seed-round-ready metrics

Proof before product.

Week 0 - Complete

Cost validation: LLM simulation confirmed ≤\$0.07/user/day at full daily usage — within budget model.

Week 1 - Complete

12 pre-paying users at \$9 upfront from a Figma mockup and 60-second demo video. Validates willingness to pay before a line of product code was written.

Now - In Progress

Backend foundation: Go + Fiber + Asynq scaffold. Gmail OAuth integrated. Claude Sonnet briefing generation running in staging with 3 beta users.

Week 5 - Target

App store submission (iOS + Android). Landing page live. Launch to 50-person waitlist built organically from the demo video.

Month 2 - Target

100 paying users · \$1.2K MRR · D7 retention >55%

MARKET DEMAND SIGNALS

OpenClaw GitHub stars (3 weeks)	145K
Fraudulent tokens exploiting hype	\$16M
Ambient agent VC investment 2025–26	\$370M
EVE pre-paying users (pre-product)	12

North star metric

Weekly Active Approvals per user — how many times a user approves an EVE-drafted reply per week. Measures genuine utility, not vanity opens. Every churned user gets a DM within 24 hours.

Technical proof points

Context Compiler producing fixed-size packages regardless of memory size. Graph traversal confirmed at <10ms for 2-hop queries. Go goroutine WebSocket server handling 1,000 concurrent test connections on a \$6/month Fly.io machine.

Built by someone who lives the problem.



Rosnel Kana
Founder & CEO
Douala, Cameroon

Fullstack engineer with 6+ years building scalable web and AI-driven systems. Former Technical Development Manager at Next Venture Tech. Currently building at the intersection of AI, mobile, and infrastructure — the exact convergence EVE requires.

Based in Douala — one of West Africa's most dynamic tech cities — Rosnel builds for mobile-first markets from the inside. That's not a marketing angle. It's a structural advantage.

- React Native
- Go
- NestJS
- TypeScript
- PostgreSQL
- AI Integration
- Web3
- System Architecture

PROVEN SHIPPING RECORD

EduFlowAI

Live

AI-powered school ERP platform. 500+ active users. End-to-end build: product design → backend → mobile app → deployment.

SecureWatch

In progress

Blockchain-based luxury watch authentication platform. Smart contract service layer, NFT history, fractional ownership — solving trust in a \$75B market.

Why this founder for this product

EVE requires Go concurrency, React Native mobile, PostgreSQL graph queries, LLM streaming, and WebRTC voice — all of which Rosnel has shipped in production. The stack is not new territory. The product is.

Actively seeking

A co-founder with GTM and growth background, ideally MENA-based. Hub71's network is a direct path to this. EVE's long-term roadmap also includes a Head of AI Research role once seed is closed.

Bootstrapped to proof. Now ready to scale.

PREVIOUS ROUNDS

- Fo

Bootstrapped — \$0 raised

EVE has been built entirely from personal resources and revenue from consulting work. No dilution to date. Full founder control. Every architectural decision made for product quality, not investor optics.
- Fi

Pre-revenue validation — 12 paying users

12 users paid \$9 upfront for a Figma mockup before product was built. This is the most capital-efficient form of market validation: real money from real people before a line of product code existed.


What bootstrapping proves

The architecture, the PRD, the cost model, and the initial traction were built with zero outside capital. The discipline required to do this — validating before building, designing for unit economics from day one — is the same discipline that will make EVE a capital-efficient investment.

CAP TABLE — CURRENT



POST-ROUND CAP TABLE (TARGET)



\$150K

Pre-seed round via SAFE. 18-month runway. Enough to ship V1, reach \$36K MRR, and raise a \$1M+ seed on real traction — not a deck.



At month 12 with 3,000 users and \$36K MRR, EVE raises a \$1M+ seed to accelerate the MENA market, add Outlook/Slack integrations, launch the Voice V2 pipeline, and hire the first two engineers.

Abu Dhabi isn't a stop. It's the launchpad.

Hub71 sits at the intersection of everything EVE needs: capital access, regulatory clarity, the Mubadala/G42 ecosystem for enterprise pilots, and a direct bridge to the fastest-growing mobile-first user base in the world — MENA's 400M+ young professionals.

Year 1 — Establish MENA as primary market

Launch Arabic language support and localized onboarding within 90 days of joining Hub71. Target UAE and Saudi Arabia as first paid markets outside Africa — high-income, mobile-first, underserved by US-centric productivity tools.

Enterprise pilots via G42 & Mubadala network

EVE's graph memory and proactive agent architecture has a natural B2B extension — a version for executive assistants managing C-suite communications. Hub71's network provides direct access to enterprise pilot partners that would take 2 years to reach cold.

Regulatory clarity for AI + voice products

UAE's AI regulatory framework is among the most progressive globally. Operating from Abu Dhabi gives EVE clear legal standing for voice data processing and AI agent capabilities across MENA — a compliance moat vs competitors operating from outside the region.

WHAT EVE NEEDS FROM HUB71

- **Co-founder matching**
MENA-based GTM / growth co-founder with regional network
- **Enterprise pilot introductions**
3-5 pilot companies from Mubadala / G42 portfolio for B2B proof
- **AI infrastructure partnerships**
G42 / G42 Cloud for Fly.io equivalent edge infrastructure in UAE
- **Seed follow-on network**
Warm introductions to regional VCs for the \$1M+ seed at month 12

What Hub71 gets from EVE

A mobile-first AI agent company building for the 2 billion users Big Tech structurally ignores — with a technical architecture that can't be replicated as a feature drop, a founder who ships, and a product that earns its right to exist from day one through real revenue, not hype.

Reach Rosnel

[linkedin.com/in/rosnelkana](https://www.linkedin.com/in/rosnelkana) · github.com/Silent-bat